
Phase-Adaptive Generation: Learning Block Budgets to Reduce Diffusion-LM Calls

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Diffusion language models can revise several tokens per forward call, but their
2 practical efficiency depends on deciding how large a span to attempt and when to
3 stop refining it. Phase-Adaptive Generation (PAG) learns this refinement budget
4 from prior block trajectories while retaining reactive delimiter boundaries. On
5 LLaDA-8B-Instruct, the best corrected PAG result uses 6.7% fewer model calls
6 on MATH-500 with a 1.33-point accuracy decrease; on GSM8K it saves 4.0%
7 with a 1.60-point decrease, while a block-size lookup remains stronger. We de-
8 velop this result through a sequence of increasingly controlled approaches: token-
9 level phase analysis, categorical/ordinal prediction from 138,901 AdaBlock tu-
10 ples, cross-model risk calibration, conservative budget ledgers, and verified macro-
11 steps. The cross-model policy reaches a 9.6% normalized call reduction but fails
12 its accuracy gate. Output-preserving policies save at most 1.24%, and batched
13 verification breaks trajectory equivalence on 22 of 64 pilot prompts despite larger
14 apparent call reductions. Together, the experiments show where learned schedul-
15 ing helps and where it must yield control: histories can allocate refinement ca-
16 pacity, but token commitment requires current-state boundaries and an audited,
17 batch-invariant reference transition.

18 1 Introduction

19 Diffusion language models generate by repeatedly revising masked positions, so one model call can
20 advance several tokens at once. This parallelism creates a control problem absent from left-to-right
21 decoding: how large a span should the model attempt, and how long should it refine that span before
22 moving on? Existing decoders answer with fixed schedules or current-step confidence signals [Wu
23 et al., 2026, Lu et al., 2026, Mohamed et al., 2025]. Both are useful, but neither asks whether the
24 trajectory itself forecasts where later compute will be needed.

25 Phase-adaptive generation (PAG) starts from a simple observation: reasoning traces are not uni-
26 formly difficult. Template text and delimiters often stabilize early, while arithmetic, variable bind-
27 ing, and answer finalization remain uncertain for longer. The difficult part is turning this structure
28 into a decoder. Offline accuracy is not enough. The scheduler changes the trajectory that produces
29 its own future features; block sizes are strongly multimodal; and a premature commitment can alter
30 every later block. A useful predictor can therefore be a poor controller.

31 We develop the method from trace analysis through deployment. We first instrument Dream and
32 LLaDA [Ye et al., 2025, Nie et al., 2025], measure token stabilization, and use change-point detec-
33 tion (CPD) to expose local shifts in uncertainty. Because those boundaries depend strongly on the
34 detector, we then move to 138,901 block-native tuples collected from AdaBlock on 5,000 GSM8K

training problems. A categorical/ordinal Transformer predicts the next block control from recent history, while the deployed scheduler preserves AdaBlock’s delimiter boundary and uses the prediction only to set a refinement budget. Each subsequent approach addresses a failure exposed by the previous one: cross-model calibration tests transfer, risk ledgers protect outputs, and verified macro-steps try to recover useful savings without granting the predictor acceptance authority.

Across the experiments, corrected forward-pass accounting shows that PAG reduces the mean number of forward evaluations (NFE) on LLaDA by 4.0% on GSM8K and 6.7% on MATH-500, with accuracy decreases of 1.60 and 1.33 percentage points. A size-conditioned lookup table is at least as competitive on GSM8K. A later cross-model policy lowers normalized NFE by 9.6% overall, but fails the predeclared accuracy-noninferiority gate because Dream loses substantially more accuracy. Conservative risk ledgers preserve outputs but save less than 1.3%; batched macro-verification appears promising in physical NFE yet fails exact trajectory audits. The same limitation recurs across these protocols: phase histories can forecast computational effort, but they should not decide token acceptance without verification.

Contributions. We make three empirical contributions. First, we connect token-level phase diagnostics to a block-native prediction problem and document why continuous boundary regression fails on a multimodal control space. Second, we give an audited evaluation of learned block-history scheduling, including a stronger history-free control and corrected NFE accounting. Third, we trace seven artifact-backed protocols across two model families, showing how risk calibration, conservative routing, and verified decoding change the efficiency–quality frontier. The outcome is a concrete rule for adaptive dLM decoding: learn capacity allocation, retain reactive semantic boundaries, and keep commitment under an auditable token-level authority.

2 Related work

Masked and block diffusion language models. Discrete and masked diffusion models replace left-to-right factorization with iterative denoising [Lou et al., 2024, Sahoo et al., 2024]. LLaDA and Dream scale this formulation to instruction following and reasoning [Nie et al., 2025, Ye et al., 2025]. Block diffusion introduces a semi-autoregressive interface: the decoder completes one contiguous block before exposing it as context to the next [Arriola et al., 2025]. That interface makes block size and within-block refinement explicit controls. Our work studies whether realized block histories help set those controls; it does not change the model or training objective.

Adaptive and early-exit decoding. Fast-dLLM combines confidence transfer with KV caching and parallel decoding [Wu et al., 2026]. AdaBlock chooses semantic boundaries from delimiter confidence and adapts refinement to the active block [Lu et al., 2026]; it is our native decoder and strongest same-stack baseline. Adaptive Parallel Decoding (APD) learns token-wise denoising paths, whereas SchED and more recent stability rules stop from inter-step similarity or persistent predictions [Israel et al., 2025, Mohamed et al., 2025, 2026]. These methods establish that current-state confidence and stability are useful. PAG asks a different question: whether previous blocks add predictive value after controlling for a strong current-state or size-conditioned rule.

Risk control and verified acceleration. Learn-then-Test and conformal risk control provide finite-sample tools for selecting a policy under a bounded loss [Angelopoulos et al., 2021, 2022]. We apply this machinery to a same-state premature-commitment loss, while evaluating task correctness separately. Self-speculative dLM decoders instead batch draft states and verify them against a reference transition [Gao et al., 2025, Han et al., 2026]. Our later protocols combine these ideas but also audit a systems assumption that is easy to omit: a nominally deterministic transition must remain invariant when evaluated alone or as one row of a batch. What remains unclear in this literature is whether a predictive phase signal survives both closed-loop deployment and the batch implementation of verification. Our contribution is the resulting end-to-end evidence boundary, not a claim of superiority over these contemporaneous decoders.




Figure 1: PAG separates two control signals. Current logits determine the next semantic boundary; history predicts a refinement budget. The realized tuple closes the loop only after the block is decoded.




Figure 2: From token-level phases to block-native data. (a) One algebra trace contains adjacent spans with markedly different stabilization confidence; the exact boundaries vary with detector settings. (b) The 138,901 training blocks form a discrete, multimodal distribution. Size 1 is dominated by delimiter blocks, while size 16 is the default content cap.

83 3 Phase-adaptive generation

84 PAG has three stages (Figure 1). Offline trace analysis tests whether the trajectory contains localized
85 compute signals. Block-native traces then supervise a next-control predictor. Online, the predictor
86 shares control with the native decoder: AdaBlock selects the semantic boundary from current log-
87 its, while PAG forecasts only the refinement budget. This division was fixed before the audited
88 evaluation.

89 3.1 Phase signals are diagnostic, not labels

90 For each output position, we record its predicted identity, top-one and top-two probabilities, and
91 entropy after every refinement call. The stabilization time of token i is the earliest step after which
92 its identity never changes,

$$s_i = \min\{t : \hat{x}_i^{(u)} = \hat{x}_i^{(T)} \text{ for all } u \geq t\}. \quad (1)$$

93 We read confidence, margin, entropy, and step index at the stabilization time in Equation 1, then
94 segment each feature along token position with Pruned Exact Linear Time (PELT) [Killick et al.,
95 2012]. Figure 2 shows the recurring pattern: long high-confidence spans are interrupted by lower-
96 confidence computation, and block sizes concentrate on delimiter-only and default-cap modes.

97 The segmentation is intentionally not used as supervision. Boundaries move when the feature,
98 smoothing window, cost, or penalty changes; representative traces yield between two and eight
99 segments under reasonable settings. CPD therefore establishes that trajectories are structured, but
100 does not define a stable scheduler target. The training path instead uses decisions directly observed
101 from a block decoder.

3.2 Block-native next-control prediction

We run AdaBlock with LLaDA-8B-Instruct on 5,000 GSM8K training problems [Cobbe et al., 2021]. After each decoded block, we store its size and realized NFE, stabilization/refinement statistics, mean and minimum confidence, digit fraction, delimiter fraction, and gap statistics. The corpus contains 138,901 training tuples and 5,543 validation tuples. Sliding windows of eight realized tuples predict the next pair $y_t = (b_t, r_t)$, where b_t is block size and r_t is the refinement budget.

The target distribution rules out ordinary regression. A squared-error predictor averages the size-1 delimiter mode and the content-block mode, producing boundaries that seldom occur in the trace corpus. We instead use a Transformer encoder [Vaswani et al., 2017] with a 128-way categorical head for b_t and an 83-threshold ordinal head for r_t . The budget head is conditioned on the block-size logits. With categorical logits z^b and ordinal logits z^r , training minimizes

$$\mathcal{L} = \text{CE}(z^b, b_t) + \frac{1}{83} \sum_{k=1}^{83} \text{BCEWithLogits}(z_k^r, \mathbf{1}[r_t > k]). \quad (2)$$

We optimize Equation 2 with a model of hidden size 64, four heads, four layers, 0.1 dropout, and learning rate 5×10^{-4} ; all inputs are z-score normalized with training statistics. Appendix B reports regression, random-forest, and architecture ablations.

3.3 Hybrid online scheduling

At block t , AdaBlock scans current logits for a high-confidence delimiter and returns the boundary b_t . The first block follows AdaBlock; later blocks pass the last eight realized tuples to the predictor and use only its budget r_t . Within the active block, masked positions are refined until the budget is reached. The decoder may then exit when all remaining predictions exceed confidence 0.8 or have stayed unchanged for two calls; a hard ceiling prevents unbounded refinement. Each initial proposal call and each refinement call counts toward NFE.

This hybrid follows directly from an observed feedback failure. Letting the predictor choose both controls often selects $(b_t, r_t) = (1, 1)$. The resulting trivial tuple makes another small prediction more likely, pushing the online feature distribution farther from the training corpus at every block. Preserving the reactive delimiter boundary breaks that loop. It also leaves a clean experimental question: does history forecast budget better than AdaBlock or a lookup table conditioned only on b_t ?

4 Evidence protocol

Initial audited evaluation. The primary PAG comparison uses deterministic LLaDA-8B-Instruct decoding on all 1,319 GSM8K test problems and a 300-example index-stratified MATH-500 sample [Hendrycks et al., 2021]. GSM8K also includes a frozen development slice for choosing among eight scheduler variants and a disjoint 919-example confirmatory complement. The comparators are AdaBlock and a history-free lookup policy that returns the median training NFE for the current block size. Every method counts the boundary proposal call and all subsequent refinement calls. This corrects the legacy course result, whose PAG counter omitted the initial proposal and is excluded from every claim here.

Cross-model evaluation. Risk-calibrated PAG (RC-PAG) uses pinned LLaDA-8B-Instruct and Dream-v0-Instruct-7B revisions, temperature zero, generation length 256, at most 64 refinement calls, and block length at most 64. Estimator fitting, policy screening, and calibration use disjoint mixtures of GSM8K, MATH, and Mostly Basic Python Problems (MBPP) training data [Austin et al., 2021]. Confirmation uses 500 GSM8K test, 300 MATH-500, 100 sanitized MBPP, and 100 HumanEval prompts [Chen et al., 2021]; HumanEval is out of domain. Each candidate stop is compared with a same-state full-refinement shadow. A finite policy family is selected with exact binomial Learn-then-Test control at familywise level 0.05, but the certificate concerns this premature-commitment loss, not answer correctness.

Promotion gates and statistics. Results are paired by prompt. We report mean NFE, exact task accuracy, and latency, with 10,000-sample paired bootstrap intervals where available. A cross-model




Figure 3: Approach progression. Each stage addresses the limitation exposed immediately before it; reported savings are relative to AdaBlock. The strongest corrected task result is PAG’s 6.7% NFE reduction on MATH-500 with a 1.33-point accuracy decrease. Larger apparent reductions fail either task-quality or exact-equivalence gates; output-preserving screens save at most 1.24%.

Table 1: Corrected LLaDA evaluation. Accuracy is exact-answer percentage (higher is better); NFE counts all model forward calls (lower is better). The lookup baseline conditions only on block size.

Dataset	Method	Accuracy $\uparrow$	Mean NFE $\downarrow$	Δ NFE
GSM8K ($n = 1,319$)	AdaBlock	77.79	87.81	—
	size lookup	76.42	82.14	−6.5%
	PAG	76.19	84.26	−4.0%
MATH-500 ($n = 300$)	AdaBlock	38.00	119.65	—
	PAG	36.67	111.67	−6.7%

policy is promoted only if it reduces NFE on both models, beats the eligible history-free control, and has an accuracy-difference lower bound of at least -2 percentage points. Later protocols add model-wise minimum savings or exact sequence-equivalence gates. Partial screens and failed audits are reported as diagnostics, never as confirmatory wins.

5 Results

Figure 3 gives the experimental logic. We begin with the strongest corrected PAG result, then change one form of control at a time. A stage advances only when its evidence passes the corresponding accuracy, savings, or equivalence gate.

5.1 PAG reduces corrected NFE by up to 6.7%

Table 1 gives the corrected LLaDA result. Relative to AdaBlock, PAG saves 4.0% NFE on GSM8K and 6.7% on MATH-500, but loses 1.60 and 1.33 accuracy points. On the untouched GSM8K complement, the paired accuracy difference is -1.20 points with a 95% interval of $[-2.61, 0.22]$. The interval includes parity, but the point estimates do not support a dominance claim.

The size lookup is the decisive control: it uses less NFE than PAG on GSM8K and is slightly more accurate. History therefore does not explain the full gain. It does help forecast heterogeneous blocks offline, but much of the deployable signal is already captured by block size. The synchronized 150-prompt timing subset points in the same cautious direction: both decoders reach 84% accuracy and PAG is 0.31 seconds faster on average, but the paired latency interval crosses zero.

5.2 Risk calibration increases savings but exposes transfer failure

The cross-model experiment sharpens the distinction between prediction and control. Histogram-gradient-boosting risk estimators are strong on prompt-grouped holdouts: history features reach area under the receiver operating characteristic curve (AUROC)/Brier scores of 0.982/0.051 for LLaDA and 0.990/0.026 for Dream. All frozen low-risk candidates also satisfy their finite-sample shadow-loss tests. Yet Table 2 shows that this surrogate success does not transfer uniformly to answer accuracy.

Across both models and all four evaluation sets, normalized NFE falls from 1.000 to 0.904, compared with 0.968 for the best eligible nonlearned rule. This aggregate is not a positive headline: the predeclared accuracy-noninferiority gate fails, with a minimum lower confidence bound of -21.2 points. LLaDA remains close to parity but saves only 1–2%; Dream obtains larger compute re-

Table 2: Cross-model confirmation for the history policy. ΔAcc is the accuracy change in percentage points and ΔNFE is relative to AdaBlock.

Model	Dataset	AdaBlock Acc./NFE	RC-PAG Acc./NFE	ΔAcc	ΔNFE
LLaDA	GSM8K	78.2 / 86.89	78.2 / 85.06	+0.0	-2.1%
LLaDA	MATH-500	37.7 / 119.72	37.3 / 118.17	-0.4	-1.3%
Dream	GSM8K	69.8 / 103.04	53.6 / 89.46	-16.2	-13.2%
Dream	MATH-500	45.3 / 123.56	40.3 / 114.39	-5.0	-7.4%

ductions by changing too many answers. A well-calibrated same-state stop label is therefore an incomplete proxy for end-to-end quality under a new decoder family.

5.3 Conservative routing converges on output preservation

The next four protocols ask whether tighter routing can keep the savings without changing answers. V4 applies a local risk threshold and saves 9.1%/3.2% NFE on LLaDA/Dream, but Dream misses the predeclared 5% floor. V5 predicts both benefit and harm; harm AUROC is only 0.372/0.456, so the router is not promoted. V6 replaces local decisions with a prompt-level risk ledger, and v7 relaxes that gate. Both record zero harmful regressions in their screens, but save only 0.75%/0.45% and 1.24%/0.65%. This is the best output-preserving point reached by learned routing, but it gives back most of PAG’s original efficiency gain. Protocols v2–v3 were superseded before GPU evidence was materialized and are not counted as experiments.

5.4 Verified macro-steps recover calls but not exactness

Under v8, a reference check evaluates a batch of fixed-depth proposals in one call. This lowers physical invocations from 83.47 to 71.19 on LLaDA and from 97.63 to 76.81 on Dream, reductions of 14.7% and 21.3%. These are the largest raw call reductions in the progression. They are not lossless results: 22 of 64 prompt trajectories disagree with AdaBlock. LLaDA latency also rises from 3.09 to 3.83 seconds because each invocation evaluates more rows. V9 audits the mechanism before another rollout. Execution fingerprints match. Successor states differ across depth-one and depth-two batch shapes. Figure 4 in the appendix shows this exactness test. The failed gate yields a final division of responsibility: learned histories propose budgets; a batch-invariant reference transition retains acceptance authority.

6 Discussion

Phase structure is useful but not sufficient. The trace analysis, tuple predictors, and risk estimators agree that refinement trajectories contain signal. The deployment results identify what that signal supports: forecasting relative effort. They do not support handing a learned score sole authority to commit a block. Current-state boundary and stability rules use contemporaneous evidence [Lu et al., 2026, Mohamed et al., 2025, 2026]; the problematic transfer here occurs when history predicts a future commitment. This difference explains why high risk-model AUROC can coexist with poor Dream accuracy and why reducing the score threshold can preserve outputs while eliminating the compute benefit.

History-free controls should be mandatory. Block size encodes both span length and AdaBlock’s semantic-boundary decision. A lookup table on that single variable beats the learned scheduler on GSM8K, so comparisons against a fixed budget or AdaBlock alone overstate the value of history. Future learned schedulers should predict a bounded residual over a strong current-state prior and show that history improves a paired frontier, rather than reporting only an absolute saving.

Verification is an implementation property. The speculative pilots suggest that batching reference transitions could reduce sequential depth, but the failed equivalence audits prevent a lossless claim. Exactness requires more than matching a draft token with a batched logit row: the reference operator, cache state, attention mask, block boundary, and batch shape must induce the same complete successor as serial decoding. This batch-invariance audit complements token-level verification

in self-speculative decoders [Gao et al., 2025, Han et al., 2026]. Physical model calls, evaluated rows, latency, and memory are separate costs and should be reported separately.

7 Limitations

The study covers two 7–8B model families, 256-token generations, and math/code benchmarks; it does not show that phase history is unhelpful for larger models, longer contexts, or a scheduler trained jointly with the model. CPD evidence is qualitative and detector-dependent. The risk certificates apply to a declared shadow loss and calibration mixture, not semantic correctness or arbitrary deployment data. Code-task exact-match rates are sensitive to extraction and test harnesses, so they appear in the appendix rather than drive the main claim. Latency was measured at batch size one on a single-GPU workflow and is not a serving benchmark. Finally, the v8–v9 results diagnose this implementation of batched verification; they do not rule out verified acceleration with stricter batch-invariant kernels or serial verification.

8 Conclusion

PAG began with the hypothesis that reasoning phases should receive different refinement budgets. The hypothesis is partly right: token stabilization is structured, block histories predict compute, and a hybrid scheduler reduces LLaDA model calls on GSM8K and MATH-500. The stronger claim does not survive the full evidence. A block-size lookup is competitive, cross-model risk control fails task-accuracy transfer, conservative gates save little, and speculative batching fails exactness audits.

Our experiments support a separation of responsibilities. Learned phase models are well suited to proposing how much capacity to allocate. Semantic boundaries should remain tied to current model evidence, and token acceptance should remain with an audited reference transition. A next system should therefore combine a bounded history residual with batch-invariant verification and treat an unsuccessful equivalence or task-quality gate as a normal fallback, not a result to average away.

References

- Anastasios N. Angelopoulos, Stephen Bates, Emmanuel J. Candès, Michael I. Jordan, and Lihua Lei. Learn then test: Calibrating predictive algorithms to achieve risk control, 2021. URL <https://arxiv.org/abs/2110.01052>.
- Anastasios N. Angelopoulos, Stephen Bates, Adam Fisch, Lihua Lei, and Tal Schuster. Conformal risk control, 2022. URL <https://arxiv.org/abs/2208.02814>.
- Marianne Arriola, Aaron Gokaslan, Justin T. Chiu, Zhihan Yang, Zhixuan Qi, Jiaqi Han, Subham Sekhar Sahoo, and Volodymyr Kuleshov. Block diffusion: Interpolating between autoregressive and diffusion language models. In *International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=tyEYt267x>.
- Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, and Charles Sutton. Program synthesis with large language models, 2021. URL <https://arxiv.org/abs/2108.07732>.
- Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, Alex Ray, Raul Puri, Gretchen Krueger, Michael Petrov, Heidy Khlaaf, Girish Sastry, Pamela Mishkin, Brooke Chan, Scott Gray, Nick Ryder, et al. Evaluating large language models trained on code, 2021. URL <https://arxiv.org/abs/2107.03374>.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems, 2021. URL <https://arxiv.org/abs/2110.14168>.
- Yifeng Gao, Ziang Ji, Yuxuan Wang, Biqing Qi, Hanlin Xu, and Linfeng Zhang. Self speculative decoding for diffusion large language models, 2025. URL <https://arxiv.org/abs/2510.04147>.

266 Ligong Han, Hao Wang, Han Gao, Kai Xu, and Akash Srivastava. S2D2: Fast decoding for diffusion
267 LLMs via training-free self-speculation, 2026. URL <https://arxiv.org/abs/2603.25702>.

268 Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song,
269 and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset, 2021.
270 URL <https://arxiv.org/abs/2103.03874>.

271 David Israel, Guy Van den Broeck, and Aditya Grover. Adaptive parallel decoding for diffusion
272 language models. In *Advances in Neural Information Processing Systems*, 2025. URL <https://openreview.net/forum?id=xwqTt26NJf>.

274 Rebecca Killick, Paul Fearnhead, and Idris A. Eckley. Optimal detection of changepoints with a
275 linear computational cost. *Journal of the American Statistical Association*, 107(500):1590–1598,
276 2012. doi: 10.1080/01621459.2012.737745.

277 Aaron Lou, Chenlin Meng, and Stefano Ermon. Discrete diffusion modeling by estimating the
278 ratios of the data distribution. In *International Conference on Machine Learning*, 2024. URL
279 <https://proceedings.mlr.press/v235/lou24a.html>.

280 Guanxi Lu, Hao Chen, Yuto Karashima, Zhican Wang, Daichi Fujiki, and Hongxiang Fan.
281 AdaBlock-dLLM: Semantic-aware diffusion LLM inference via adaptive block size. In *International
282 Conference on Learning Representations*, 2026. URL <https://arxiv.org/abs/2509.26432>.

284 Amr Mohamed, Yang Zhang, Michalis Vazirgiannis, and Guokan Shang. SchED: Similarity-based
285 early stopping for efficient diffusion language model decoding, 2025. URL <https://arxiv.org/abs/2512.02892>.

287 Amr Mohamed, Guokan Shang, and Michalis Vazirgiannis. LESS is more: Mutual-stability sam-
288 pling for diffusion language models, 2026. URL <https://arxiv.org/abs/2606.16908>.

289 Shen Nie, Fengqi Zhu, Zebin You, Xiaolu Zhang, Jingyang Ou, Jun Hu, Jun Zhou, Yankai Lin, Ji-
290 Rong Wen, and Chongxuan Li. Large language diffusion models, 2025. URL <https://arxiv.org/abs/2502.09992>.

292 Subham Sekhar Sahoo, Marianne Arriola, Yair Schiff, Aaron Gokaslan, Edgar Marroquin, Justin T.
293 Chiu, Alexander M. Rush, and Volodymyr Kuleshov. Simple and effective masked diffusion
294 language models. In *Advances in Neural Information Processing Systems*, volume 37, 2024. doi:
295 10.52202/079017-4135.

296 Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez,
297 Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural
298 Information Processing Systems*, 2017. URL <https://proceedings.neurips.cc/paper/7181-attention-is-all-you-need>.

300 Chengyue Wu, Hao Zhang, Shuchen Xue, Zhijian Liu, Shizhe Diao, Ligeng Zhu, Ping Luo, Song
301 Han, and Enze Xie. Fast-dLLM: Training-free acceleration of diffusion LLM by enabling KV
302 cache and parallel decoding. In *International Conference on Learning Representations*, 2026.
303 URL <https://arxiv.org/abs/2505.22618>.

304 Jiacheng Ye, Zhihui Xie, Lin Zheng, Jiahui Gao, Zirui Wu, Xin Jiang, Zhenguo Li, and Lingpeng
305 Kong. Dream 7b: Diffusion large language models, 2025. URL <https://arxiv.org/abs/2508.15487>.

307 A Verification equivalence audit




Figure 4: Exactness audit for verified macro-steps. A batched row is eligible only when it induces the same complete successor as the serial reference transition. The v8 pilot disagrees on 22 of 64 prompts, and v9 rejects the depth-dependent batch shape before rollout.

309 B Predictor and integration diagnostics

310 **Continuous regression.** The first Transformer predicts real-valued controls with normalized val-
 311 idation loss 0.096. Deployed on 30 arithmetic prompts, it reaches 73% accuracy versus 93% for
 312 AdaBlock and averages 8.1 blocks rather than 12.3. Inspection shows interpolated block sizes merg-
 313 ing semantic spans and producing repeated or truncated text. This run motivated the categorical
 314 block head; it is not used as a baseline in the audited evaluation.

315 **Random forest.** A 200-tree, depth-15 multi-output forest uses rolling statistics from a five-tuple
 316 window. On 90,994 training and 22,907 held-out examples, block-size mean absolute error (MAE)
 317 is 7.24 and stabilization-step MAE is 1.66. The last block size is the largest individual feature
 318 (importance 0.214), supporting the explicit size-lookup control.

319 **Transformer ablation.** The categorical/ordinal sweep varies context length, hidden size, attention
 320 heads, layers, dropout, and learning rate. The selected $(w, d, h, l) = (8, 64, 4, 4)$ model reaches
 321 validation loss 2.216957. Wider models do not improve this objective: at four layers, the best losses
 322 for widths 32, 64, 128, 256, and 512 are 2.233, 2.217, 2.238, 2.289, and 2.336.

323 C Complete cross-model confirmation

324 Table 3 reports every dataset/model pair from the v1 confirmation; the code tasks are retained for
 325 completeness but do not drive the main claim.

Table 3: Complete v1 confirmation. Accuracy is percentage and NFE is the mean number of physical model calls. HumanEval is declared out of domain.

Model	Dataset	Method	Accuracy	Mean NFE
LLaDA	GSM8K	AdaBlock / RC-PAG	78.2 / 78.2	86.89 / 85.06
LLaDA	MATH-500	AdaBlock / RC-PAG	37.7 / 37.3	119.72 / 118.17
LLaDA	MBPP	AdaBlock / RC-PAG	4.0 / 4.0	71.66 / 70.57
LLaDA	HumanEval	AdaBlock / RC-PAG	32.0 / 33.0	95.48 / 94.39
Dream	GSM8K	AdaBlock / RC-PAG	69.8 / 53.6	103.04 / 89.46
Dream	MATH-500	AdaBlock / RC-PAG	45.3 / 40.3	123.56 / 114.39
Dream	MBPP	AdaBlock / RC-PAG	3.0 / 0.0	80.02 / 54.25
Dream	HumanEval	AdaBlock / RC-PAG	48.0 / 16.0	84.78 / 64.50

326 D Artifact-backed protocol ledger

327 Table 4 records the evidence source for each reported phase. Hashes are shortened only for readabil-
 328 ity. All records are non-mock and pin model and dataset revisions in their environment manifests.

Table 4: Artifact provenance for numerical claims.

Evidence	Config/run hash	Authoritative artifact
Initial PAG	Strategy 1	regraded report summarized in final report
RC-PAG v1	5197d981c0bf	report/summary.json, claim_audit.json
Local risk v4	8cce1af8e35b	screening_summary.json
Advantage v5	7688c5235bd4	screen records, advantage_manifest.json
Risk ledger v6	58e76c3eb9b1	readiness_audit.json
Relaxed gate v7	c1eda289fb08	readiness_audit.json
Verified pilot v8	d36b982c2388	parity_audit.json
Equivalence audit v9	588204cfb482	equivalence_audit.json

E Accounting and reporting rules

NFE denotes physical invocations of the model forward function, including the initial proposal used to choose a boundary. A batched invocation counts as one physical NFE and separately contributes one evaluated row per batch element. We do not translate NFE reductions into FLOP, energy, or latency claims. Accuracy is paired on identical prompt IDs. A run with a failed manifest, incomplete coverage, a mock record, or a failed claim gate is ineligible for a positive aggregate.

F Reproducibility, compute, and assets

Configurations and statistics. The anonymous supplement contains the implementation, exact YAML configurations, pinned revisions, run manifests, and artifact summaries referenced in Table 4. The initial audit uses seed 20260710; the cross-model protocols use seed 20260729. Confidence intervals are percentile intervals from 10,000 paired bootstrap resamples of prompt IDs. The risk tests use exact binomial bounds at familywise level 0.05. The supplement documents the commands for the initial Strategy 1 evaluation and each RC-PAG protocol. Raw checkpoints and benchmark copies are not redistributed; the loaders fetch the pinned public assets.

Compute. We estimate approximately 160 A100-equivalent GPU-hours for the completed experiments. Trace collection and the original PAG study account for about 40 hours; the six materialized protocols v1 and v4–v8 account for about 20 hours each, or 120 hours. V2–v3 were superseded before full GPU runs, and v9 stopped at the equivalence preflight before a benchmark rollout; incomplete job attempts and CPU-only analysis are excluded. The estimate assumes one NVIDIA A100 PCIe (80 GB), with LLaDA and Dream loaded one at a time in bfloat16. The recorded environment used Python 3.11.15, PyTorch 2.5.1 with CUDA 12.1, Transformers 4.46.3, Datasets 4.8.5, scikit-learn 1.9.0, and math-verify 0.9.0. Atomic records retain per-prompt elapsed time; the v8 pilots ranged from roughly one to six seconds per method/prompt.

Licenses. The pinned LLaDA checkpoint is MIT licensed and the pinned Dream checkpoint is Apache-2.0 licensed. GSM8K, MATH-500, and HumanEval are used under their MIT licenses; MBPP data are used under CC BY 4.0. We cite each model and benchmark at first use and preserve their original terms. The supplement contains no model weights or benchmark examples beyond the aggregate statistics and one short diagnostic trace shown in Figure 2.

G Broader impact

Reducing sequential model calls can lower latency and energy when the extra batch work does not erase the gain. It can also make generative systems cheaper to deploy at scale. The experiments expose the corresponding reliability risk: an efficiency controller that commits tokens too early can retain attractive proxy metrics while changing answers substantially. Systems using learned scheduling should therefore keep a reference fallback, audit task quality, and report batch work as well as call count. This work releases no new foundation model or scraped dataset and does not expand the underlying models’ capabilities; any downstream misuse or bias remains that of the base models and their deployment context.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and Section 1 state the evaluated settings, corrected efficiency results, accuracy trade-offs, failed gates, and the resulting scope of the design claim. Sections 5 and 7 support and delimit those claims.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 7 discusses model scale, benchmark scope, detector dependence, calibration scope, latency, and the limits of the failed verification audits.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper contains no theorem or formal proof. The finite-sample risk-control procedure invokes established methods and states the loss, policy family, and confidence level in Section 4.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes]

Justification: Sections 3 and 4 specify the scheduler, features, model settings, splits, controls, accounting, gates, and statistics. The appendix and anonymous supplement provide pinned configurations, commands, manifests, and artifact provenance.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in

some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: An anonymized source supplement accompanies the submission with the implementation, exact YAML configurations, run commands, and non-mock artifact summaries. Public model and dataset assets are fetched at the pinned revisions listed in the supplement.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes]

Justification: Architecture and training details appear in Section 3.2; decoding, dataset splits, seeds, evaluation sizes, thresholds, and policy-selection rules appear in Section 4 and the reproducibility appendix.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes]

Justification: Section 4 defines the paired design and 10,000-sample bootstrap procedure. Section 5 reports the main paired accuracy interval and the minimum confirmatory lower bound; null or failed intervals are not presented as wins.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes]

Justification: The compute appendix reports a total estimate of approximately 160 A100-equivalent GPU-hours: 40 hours for trace collection and the original PAG experiments, plus six completed artifact-backed protocols at approximately 20 hours each. It also reports GPU memory, execution mode, software versions, exclusions, and representative per-prompt times.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes]

Justification: The work evaluates public models on public benchmarks, preserves double-blind anonymity, reports failed promotion gates, and does not involve private data, human subjects, or a high-risk model release.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes]

Justification: The Broader Impact appendix discusses possible efficiency and energy benefits, lower deployment cost, reliability failures from premature commitment, scaling of underlying model harms, and reference-fallback mitigation.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.
- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A]

Justification: The paper releases no new pretrained model, scraped dataset, or capability-enhancing checkpoint. The source supplement contains only decoding and evaluation code, configurations, and aggregate artifacts.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes]

Justification: Original model and benchmark creators are cited at first use. The reproducibility appendix states the pinned model licenses (MIT and Apache-2.0) and benchmark licenses (MIT and CC BY 4.0); the supplement does not redistribute checkpoints or datasets.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [Yes]

Justification: The new scheduler implementation, configurations, and evaluation artifacts are documented in the anonymous supplement, including exact commands, schema-checked manifests, run hashes, and known limitations.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A]

Justification: The research does not use crowdsourcing or human participants.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional review board (IRB) approvals or equivalent for research with human subjects**

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A]

Justification: The research does not use crowdsourcing or human participants, so IRB review is not applicable.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. **Declaration of LLM usage**

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [N/A]

Justification: Generative AI tools were used only for writing, editing, and formatting. The research questions, method design, implementation, experiments, evaluation, analysis, and scientific interpretation were performed by the authors; no LLM generated labels or replaced experimental evaluation.

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.